

K - Means Clustering.

How did it do that ?

STEP 1: Choose the number K of clusters



STEP 2: Select at random K points, the centroids (not necessarily from your dataset)



STEP 3: Assign each data point to the closest centroid → That forms K clusters



STEP 4: Compute and place the new centroid of each cluster




STEP 5: Reassign each data point to the new closest centroid.
If any reassignment took place, go to STEP 4, otherwise go to FIN.



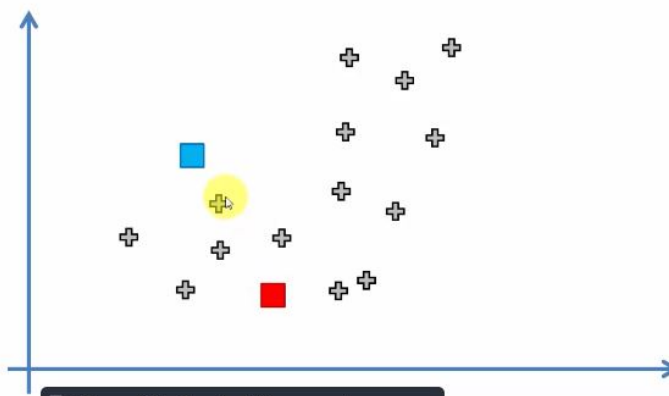
Your Model is Ready

Step by Step.


STEP 1: Choose the number K of clusters: $K = 2$




STEP 3: Assign each data point to the closest centroid → That forms K clusters




STEP 3: Assign each data point to the closest centroid → That forms K clusters



STEP 4: Compute and place the new centroid of each cluster



STEP 5: Reassign each data point to the new closest centroid.
If any reassignment took place, go to STEP 4, otherwise go to FIN.




We go back to step 4 and reassign our centroids iterate over steps 4 and 5 long enough to converge the initial and final centroids into a significantly same position.

Choosing the correct number of clusters i.e the value of K.

Within Clusters Sum of Squares:

$$WCSS = \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster 3}} \text{distance}(P_i, C_3)^2$$



Monotonically decreasing curve, choose the optimal number of clusters using the Elbow method, wherein there is no sharp decrease in the value of WCSS.

Importing the libraries.
Getting the dataset.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

os.chdir(r'C:\Users\acer\Desktop\P14-Machine-Learning-AZ-Template-Folder\Machine Learning A-Z Template Fold
df = pd.read_csv('Mall_Customers.csv')
df.head()
```

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [3]: df.columns.values

array(['CustomerID', 'Genre', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)'], dtype=object)
```

```
In [6]: df.isnull().sum()
```

```
CustomerID      0
Genre           0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

```
In [80]: df.shape
```

```
(200, 7)
```

Labelling the age groups for visualization.

```
In [17]: def ageCoder(myage):  
         if myage>=18 and myage<=30:  
             return "Youth"  
         elif myage>30 and myage<=60:  
             return "Working"  
         elif myage>60:  
             return "Senior"
```

Syntax:

```
s.apply(func, convert_dtype=True, args=())
```

Parameters:

func: .apply takes a function and applies it to all values of pandas series.

convert_dtype: Convert dtype as per the function's operation.

args=(): Additional arguments to pass to function instead of series.

Return Type: Pandas Series after applied function/operation.

```
In [22]: df['AgeClass'] = df['Age'].apply(ageCoder, convert_dtype = True)
df.head(3)
```

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)	AgeClass
0	1	Male	19	15	39	Youth
1	2	Male	21	15	81	Youth
2	3	Female	20	16	6	Youth

Labelling Annual Income for visualizaion.

```
In [51]: range = income.max() - income.min()
lim1 = income.min() + (1/3)*range
lim2 = income.max() - (1/3)*range

print(income.min(), round(lim1), round(lim2), income.max())
```

```
15 56.0 96.0 137
```

```
In [52]: def incomeCoder(salary):
    if salary<= lim1:
        return 'low'
    elif salary>lim1 and salary<=lim2:
        return 'medium'
    elif salary>lim2:
        return 'high'
```




```
In [53]: df['IncomeClass'] = df['Annual Income (k$)'].apply(incomeCoder, convert_dtype = 1)
df.head(2)
```

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)	AgeClass	IncomeClass
0	1	Male	19	15	39	Youth	low
1	2	Male	21	15	81	Youth	low

Data Visualization.

```
In [4]: plt.figure(figsize = (10,8))  
sns.set(style = 'whitegrid', font_scale = 1.2)  
  
sns.countplot(df.Genre)  
plt.show()
```



More number of girls than boys.


```
In [5]: df.Genre.value_counts()
```

```
Female    112  
Male       88  
Name: Genre, dtype: int64
```

```
In [10]: plt.figure(figsize = (15,2))
sns.set(style = 'whitegrid', font_scale = 1.2)

sns.swarmplot(df.Age)
sns.boxplot(df.Age, color = 'lightgreen')

plt.show()
```




```
In [24]: pd.crosstab(df.Genre, df.AgeClass, margins = True)
```

AgeClass	Senior	Working	Youth	All
Genre				
Female	6	72	34	112
Male	11	49	28	88
All	17	121	62	200

```
In [34]: pd.crosstab(df.Genre, df.AgeClass, margins = True, normalize = 'index').round(3) * 100
```

AgeClass	Senior	Working	Youth
Genre			
Female	5.4	64.3	30.4
Male	12.5	55.7	31.8
All	8.5	60.5	31.0

```
In [25]: plt.figure(figsize = (10,8))  
sns.set(style = 'whitegrid', font_scale = 1.2)  
  
sns.countplot(df.AgeClass, hue = df.Genre)  
plt.show()
```




More number of working class females.
Minimum number of senior women.

```
In [57]: income = df.iloc[:,3].values
plt.figure(figsize = (15,2))
sns.set(style = 'whitegrid', font_scale = 1.2)



sns.swarmplot(income)
sns.boxplot(income, color = 'lightyellow')

plt.show()
```




```
In [67]: plt.figure(figsize = (15,8))
sns.set(style = 'whitegrid', font_scale = 1.2)


plt.subplot(1,2,1)
sns.countplot(df.IncomeClass, palette = 'Set2')
plt.subplot(1,2,2)
sns.countplot(df.IncomeClass, hue = df.Genre, palette = ['green', 'tomato'])
plt.show()
```




```
In [72]: spend = df.iloc[:,4].values
plt.figure(figsize = (15,2))
sns.set(style = 'whitegrid', font_scale = 1.2)

sns.swarmplot(spend, color = 'green')
sns.boxplot(spend, color = 'purple')

plt.show()
```



```
In [76]: plt.figure(figsize = (10,8))  
sns.set(style = 'whitegrid', font_scale = 1.2)  
sns.scatterplot(x = income, y = spend, marker = 'X', color = 'red')  
plt.xlabel('Annual Income (X 1000)')  
plt.ylabel('Spend Score')  
plt.show()
```



Setting up variables.

```
In [78]: x = df.iloc[:, [3,4]].values
```

Using the elbow method to find the optimal number of clusters.

```
In [93]: from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters = 3, init = 'k-means++', random_state = 42) #n_iter = 300 (default)
y_kmeans = kmeans.fit_predict(X)
```

```
"wcss.append(kmeans.inertia_)\n\nplt.plot(range(1, 11), wcss)\nplt.title('The Elbow Method')\nplt.xlabel('Number of clusters')\nplt.ylabel('WCSS')\nplt.show()"
```

```
In [94]: y_kmeans
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 1, 2, 1, 2, 1, 2, 1, 2,
       1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
       1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
       1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
       1, 2])
```

```
In [101]: X[y_kmeans == 1]
```

```
array([[ 70, 29],
       [ 71, 35],
       [ 71, 11],
       [ 71,  9],
       [ 72, 34],
       [ 73,  5],
       [ 73,  7],
       [ 74, 10],
       [ 75,  5],
       [ 76, 40],
       [ 77, 12],
       [ 77, 36],
       [ 78, 22],
       [ 78, 17],
       [ 78, 20],
       [ 78, 16],
       [ 78,  1],
       [ 78,  1],
       [ 79, 35],
       [ 81,  5],
       [ 85, 26],
       [ 86, 20],
       [ 87, 27],
       [ 87, 13]
```

```
In [104]: print(X[y_kmeans == 1, 0])
          print(X[y_kmeans == 1, 1])
```


```
[ 70  71  71  71  72  73  73  74  75  76  77  77  78  78  78  78  78
  79  81  85  86  87  87  87  88  88  93  97  98  99 101 103 103 113 120
 126 137]
[29 35 11  9 34  5  7 10  5 40 12 36 22 17 20 16  1  1 35  5 26 20 27 13
 10 13 15 14 32 15 39 24 17 23  8 16 28 18]
```

```
In [113]: print(kmeans.cluster_centers_)  # For 5 clusters.
```

```
[[55.2962963  49.51851852]
 [88.2        17.11428571]
 [26.30434783 20.91304348]
 [25.72727273 79.36363636]
 [86.53846154 82.12820513]]
```

Visualising the clusters.

```
In [96]: plt.figure(figsize = (10,8))
sns.set(style = 'whitegrid', font_scale = 1.2)
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 50, c = 'purple', label = 'Cluster 3')
plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s = 300, c = 'green', marker = 'X')
plt.title('Clusters of customers')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
plt.show()
```



Using upto 5 clusters.


```
In [109]: plt.figure(figsize = (20,15))
sns.set(style = 'whitegrid', font_scale = 1.2)

kmeans = KMeans(n_clusters = 2, init = 'k-means++', random_state = 42) #n_iter = 300 (default)
y_kmeans = kmeans.fit_predict(X)

plt.subplot(2,2,1)
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s = 300, c = 'black',
            marker = 'X', label = 'Centroids')
plt.title('Clusters of customers')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()

kmeans = KMeans(n_clusters = 3, init = 'k-means++', random_state = 42) #n_iter = 300 (default)
y_kmeans = kmeans.fit_predict(X)

plt.subplot(2,2,2)
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 50, c = 'green', label = 'Cluster 3')

plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s = 300, c = 'black',
            marker = 'X', label = 'Centroids')
plt.title('Clusters of customers')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
```

```
kmeans = KMeans(n_clusters = 4, init = 'k-means++', random_state = 42) #n_iter = 300 (default)
y_kmeans = kmeans.fit_predict(X)
```

```
plt.subplot(2,2,3)
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 50, c = 'green', label = 'Cluster 3')
plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], s = 50, c = 'purple', label = 'Cluster 4')





plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s = 300, c = 'black',
            marker = 'X', label = 'Centroids')
plt.title('Clusters of customers')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
```

```
kmeans = KMeans(n_clusters = 5, init = 'k-means++', random_state = 42) #n_iter = 300 (default)
y_kmeans = kmeans.fit_predict(X)
```

```
plt.subplot(2,2,4)
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 50, c = 'green', label = 'Cluster 3')
plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], s = 50, c = 'purple', label = 'Cluster 4')
plt.scatter(X[y_kmeans == 4, 0], X[y_kmeans == 4, 1], s = 50, c = '#af76b2', label = 'Cluster 5')

plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s = 300, c = 'black',
            marker = 'X', label = 'Centroids')
plt.title('Clusters of customers')
```

```
plt.xlabel('Annual Income (k$)')  
plt.ylabel('Spending Score (1-100)')  
plt.legend()  
plt.show()
```



Using Within Clusters Sum of Squares.


```
In [123]: wcss = []
          hh = np.arange(1,11)
          for ii in hh:
              kmeans = KMeans(n_clusters = ii, init = 'k-means++', random_state = 42)
              kmeans.fit(X)
              wcss.append(kmeans.inertia_)

          print(wcss)

[269981.28, 181363.59595959596, 106348.37306211118, 73679.78903948834, 44448.45544793371, 37233.81451071001, 30259.65720728547, 25011.8393
4915659, 21850.165282585633, 19672.07284901432]
```

Visualizing using a plot.

```
In [130]: plt.figure(figsize = (10,8))  
sns.set(style = 'whitegrid', font_scale = 1.2)  
plt.plot(hh, wcss, 'r--')  
plt.scatter(hh, wcss, s= 100,marker = 'X')  
plt.xlabel('K - values')  
plt.ylabel('Within Clusters Sum of Squares.')  
plt.show()
```




Using the elbow method, we choose $k = 5$ as our optimal value of K .

```
In [132]: kmeans = KMeans(n_clusters = 5, init = 'k-means++', random_state = 42) #n_iter = 300 (default)
y_kmeans = kmeans.fit_predict(X)

plt.figure(figsize = (10,8))
sns.set(style = 'whitegrid', font_scale = 1.2)

plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 50, c = 'green', label = 'Cluster 3')
plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], s = 50, c = 'purple', label = 'Cluster 4')
plt.scatter(X[y_kmeans == 4, 0], X[y_kmeans == 4, 1], s = 50, c = '#af76b2', label = 'Cluster 5')

plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1], s = 300, c = 'black',
            marker = 'X', label = 'Centroids')
plt.title('Clusters of customers')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
plt.show()
```

The End.

